

C.M. Bishop's PRML: Chapter 5; Neural Networks

Vasil Khalidov & Miles Hansard

Introduction

The aim is, as before, to find useful decompositions of the **target** variable;

$$t(\mathbf{x}) = y(\mathbf{x}, \mathbf{w}) + \epsilon(\mathbf{x}) \quad (3.7)$$

- ▶ $t(\mathbf{x}_n)$ and \mathbf{x}_n are the **observations**, $n = 1, \dots, N$.
- ▶ $\epsilon(\mathbf{x})$ is the **residual** error.

Linear Models

For example, recall the (Generalized) Linear Model:

$$y(\mathbf{x}, \mathbf{w}) = f \left(\sum_{j=0}^M w_j \phi_j(\mathbf{x}) \right) \quad (5.1)$$

- ▶ $\phi = (\phi_0, \dots, \phi_M)^\top$ is the fixed model **basis**.
- ▶ $\mathbf{w} = (w_0, \dots, w_M)^\top$ are the model **coefficients**.
- ▶ For regression: $f(\cdot)$ is the identity.
- ▶ For classification: $f(\cdot)$ maps to a posterior probability.

Feed-Forward Networks

Feed-forward Neural Networks generalize the linear model

$$y(\mathbf{x}, \mathbf{w}) = f \left(\sum_{j=0}^M w_j \phi_j(\mathbf{x}) \right) \quad (5.1 \text{ again})$$

- ▶ The basis itself, as well as the coefficients w_j , will be adapted.
- ▶ Roughly: the principle of (5.1) will be used twice; once to define the basis, and once to obtain the output.

Activations

Construct M linear combinations of the inputs x_1, \dots, x_D :

$$a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i \quad (5.2)$$

- ▶ a_j are the **activations**, $j = 1, \dots, M$.
- ▶ $w_{ji}^{(1)}$ are the layer one **weights**, $i = 1 \dots D$.
- ▶ $w_{j0}^{(1)}$ are the layer one **biases**.

Each linear combination a_j is transformed by a (nonlinear, differentiable) **activation function**:

$$z_j = h(a_j) \quad (5.3)$$

Output Activations

The hidden outputs $z_j = h(a_j)$ are linearly combined in layer two:

$$a_k = \sum_{j=0}^M w_{kj}^{(2)} z_j \quad (5.4)$$

- ▶ a_k are the **output activations**, $k = 1, \dots, K$.
- ▶ $w_{kj}^{(2)}$ are the layer two **weights**, $j = 1 \dots D$.
- ▶ $w_{k0}^{(2)}$ are the layer two **biases**.

The output activations a_k are transformed by the **output activation function**:

$$y_k = \sigma(a_k) \quad (5.5)$$

- ▶ y_k are the final outputs.
- ▶ $\sigma(\cdot)$ is, like $h(\cdot)$, a sigmoidal function.

The Complete Two-Layer Model

The model $y_k = \sigma(a_k)$ is, after substituting the definitions of a_j and a_k :

$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right) \quad (5.9)$$

a_j
 a_k

- ▶ $h(\cdot)$ and $\sigma(\cdot)$ are sigmoidal functions, e.g. the logistic function.

$$s(a) = \frac{1}{1 + \exp(-a)} \quad s(a) \in [0, 1]$$

- ▶ If $\sigma(\cdot)$ is the identity, then a regression model is obtained.
- ▶ Evaluation of (5.9) is called **forward propagation**.

Network Diagram

The approximation process can be represented by a network:




Figure: 5.1

- ▶ Nodes are **input**, **hidden** and **output** units. Links are corresponding weights.
- ▶ Information propagates 'forwards' from the explanatory variable \mathbf{x} to the estimated response $y_k(\mathbf{x}, \mathbf{w})$.

Properties & Generalizations

- ▶ Typically $K \leq D \leq M$, which means that the network is redundant if all $h(\cdot)$ are linear.
- ▶ There may be more than one layer of hidden units. *cf ResNet skip from*
- ▶ Individual units need not be fully connected to the next layer.
- ▶ Individual links may skip over one or more subsequent layers.
- ▶ Networks with two (cf. 5.9) or more layers are **universal approximators**.
- ▶ Any continuous function can be uniformly approximated to arbitrary accuracy, given enough hidden units. *! 2!!*
- ▶ This is true for many definitions of $h(\cdot)$, but excluding polynomials.
- ▶ There may be **symmetries** in the weight space, meaning that different choices of \mathbf{w} may define the same mapping from input to output.

Maximum Likelihood Parameters

The aim is to minimize the residual error between $\mathbf{y}(\mathbf{x}_n, \mathbf{w})$ and \mathbf{t}_n . Suppose that the target is a scalar-valued function, which is Normally distributed around the estimate:

$$p(t|\mathbf{x}, \mathbf{w}) = \mathcal{N}(t | y(\mathbf{x}, \mathbf{w}), \beta^{-1}) \quad (5.12)$$

Then it will be appropriate to consider the sum of squared-errors

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left(y(\mathbf{x}_n, \mathbf{w}) - t_n \right)^2 \quad (5.14)$$

The **maximum-likelihood** estimate of \mathbf{w} can be obtained by (numerical) minimization:

$$\mathbf{w}_{\text{ML}} = \min_{\mathbf{w}} E(\mathbf{w})$$

Maximum Likelihood Precision

Having obtained the ML parameter estimate \mathbf{w}_{ML} , the precision, β can also be estimated. E.g. if the N observations are IID, then their joint probability is

$$p\left(\{t_1, \dots, t_N\} \middle| \{\mathbf{x}_1, \dots, \mathbf{x}_N\}, \mathbf{w}, \beta\right) = \prod_{n=1}^N p(t_n | \mathbf{x}_n, \mathbf{w}, \beta)$$

The negative log-likelihood, in this case, is

$$-\log p = \beta E(\mathbf{w}_{\text{ML}}) - \frac{N}{2} \log \beta + \frac{N}{2} \log 2\pi \quad (5.13)$$

The derivative $d/d\beta$ is $E(\mathbf{w}_{\text{ML}}) - \frac{N}{2\beta}$ and so

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} 2E(\mathbf{w}_{\text{ML}}) \quad (5.15)$$

And $1/\beta_{\text{ML}} = \frac{1}{NK} 2E(\mathbf{w}_{\text{ML}})$ for K target variables.

Error Surface

The residual error $E(\mathbf{w})$ can be visualized as a surface in the weight-space:




Figure: 5.5

- ▶ The error will, in practice, be highly nonlinear, with many minima, maxima and saddle-points.
- ▶ There will be **inequivalent** minima, determined by the particular data and model, as well as **equivalent** minima, corresponding to weight-space symmetries.

Parameter Optimization

Iterative search for a local minimum of the error:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)} \quad (5.27)$$

- ▶ ∇E will be zero at a minimum of the error.
- ▶ τ is the **time-step**.
- ▶ $\Delta \mathbf{w}^{(\tau)}$ is the weight-vector **update**.
- ▶ The definition of the update depends on the choice of algorithm.

Local Quadratic Approximation

The truncated Taylor expansion of $E(\mathbf{w})$ around a weight-point $\hat{\mathbf{w}}$ is

$$E(\mathbf{w}) \simeq E(\hat{\mathbf{w}}) + (\mathbf{w} - \hat{\mathbf{w}})^\top \mathbf{b} + \frac{1}{2}(\mathbf{w} - \hat{\mathbf{w}})^\top \mathbf{H}(\mathbf{w} - \hat{\mathbf{w}}) \quad (5.28)$$

- ▶ $\mathbf{b} = \nabla E|_{\mathbf{w}=\hat{\mathbf{w}}}$ is the **gradient** at $\hat{\mathbf{w}}$.
- ▶ $(\mathbf{H})_{ij} = \frac{\partial E}{\partial w_i \partial w_j} \Big|_{\mathbf{w}=\hat{\mathbf{w}}}$ is the **Hessian** $\nabla \nabla E$ at $\hat{\mathbf{w}}$.

The gradient of E can be approximated by the gradient of the quadratic model (5.28); if $\mathbf{w} \simeq \hat{\mathbf{w}}$ then

$$\nabla E \simeq \mathbf{b} + \mathbf{H}(\mathbf{w} - \hat{\mathbf{w}}) \quad (5.31)$$

where $\frac{1}{2} ((\mathbf{H} + \mathbf{H}^\top) \mathbf{w} - \mathbf{H} \hat{\mathbf{w}} - \mathbf{H}^\top \hat{\mathbf{w}}) = \mathbf{H}(\mathbf{w} - \hat{\mathbf{w}})$, as $\mathbf{H}^\top = \mathbf{H}$.

Approximation at a Minimum

Suppose that \mathbf{w}^* is at a minimum of E , so $\nabla E|_{\mathbf{w}=\mathbf{w}^*}$ is zero, and

$$E(\mathbf{w}) = E(\mathbf{w}^*) + \frac{1}{2}(\mathbf{w} - \mathbf{w}^*)^\top \mathbf{H}(\mathbf{w} - \mathbf{w}^*) \quad (5.32)$$

- ▶ $\mathbf{H} = \nabla \nabla E|_{\mathbf{w}=\mathbf{w}^*}$ is the Hessian.
- ▶ The eigenvectors $\mathbf{H}\mathbf{u}_i = \lambda_i \mathbf{u}_i$ are **orthonormal**.
- ▶ $(\mathbf{w} - \mathbf{w}^*)$ can be represented in \mathbf{H} -coordinates as $\sum_i \alpha_i \mathbf{u}_i$.

Hence the second term of (5.32) can be written

$$\frac{1}{2}(\mathbf{w} - \mathbf{w}^*)^\top \mathbf{H}(\mathbf{w} - \mathbf{w}^*) = \frac{1}{2} \left(\sum_i \lambda_i \alpha_i \mathbf{u}_i \right)^\top \left(\sum_j \alpha_j \mathbf{u}_j \right)$$

So that

$$E(\mathbf{w}) = E(\mathbf{w}^*) + \frac{1}{2} \sum_i \lambda_i \alpha_i^2 \quad (5.36)$$

Characterization of a Minimum

The eigenvalues λ_i of \mathbf{H} characterize the stationary point \mathbf{w}^* .

- ▶ If all $\lambda_i > 0$, then \mathbf{H} is **positive definite** ($\mathbf{v}^\top \mathbf{H} \mathbf{v} > 0$).
- ▶ This is analogous to the scalar condition $\left. \frac{\partial^2 E}{\partial w^2} \right|_{w^*} > 0$.
- ▶ Zero gradient and positive principle curvatures mean that $E(\mathbf{w}^*)$ is a **minimum**.




Figure: 5.6

Gradient Descent

The simplest approach is to update \mathbf{w} by a displacement in the negative gradient direction.

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w}^{(\tau)}) \quad (5.41)$$

- ▶ This is a **steepest descent** algorithm.
- ▶ η is the **learning rate**.
- ▶ This is a **batch** method, as evaluation of ∇E involves the entire data set.
- ▶ Conjugate gradient or quasi-Newton methods may, in practice, be preferred.
- ▶ A range of starting points $\{\mathbf{w}^{(0)}\}$ may be needed, in order to find a satisfactory minimum.

Optimization Scheme

An efficient method for the evaluation of $\nabla E(\mathbf{w})$ is needed.

- ▶ Each iteration of the descent algorithm has two stages:
- ▶ I. Evaluate derivatives of error with respect to weights (involving **backpropagation** of error through the network).
- ▶ II. Use derivatives to compute adjustments of the weights (e.g. steepest descent).

Backpropagation is a general principle, which can be applied to many types of network and error function.

Simple Backpropagation

The error function is, typically, a sum over the data points $E(\mathbf{w}) = \sum_{n=1}^N E_n(\mathbf{w})$. For example, consider a linear model

$$y_k = \sum_i w_{ki} x_i \quad (5.45)$$

The error function, for an **individual** input \mathbf{x}_n , is

$$E_n = \frac{1}{2} \sum_k (y_{nk} - t_{nk})^2, \quad \text{where} \quad y_{nk} = y_k(\mathbf{x}_n, \mathbf{w}). \quad (5.46)$$

The gradient with respect to a **weight** w_{ji} is

$$\frac{\partial E_n}{\partial w_{ji}} = (y_{nj} - t_{nj}) x_{ni} \quad (5.47)$$

- ▶ w_{ji} is a particular **link** (x_i to y_j).
- ▶ x_{ni} is the **input** to the link (i -th component of \mathbf{x}_n).
- ▶ $(y_{nj} - t_{nj})$ is the **error** output by the link.

General Backpropagation

Recall that, in general, each unit computes a weighted sum:

$$a_j = \sum_i w_{ji} z_i \quad \text{with activation} \quad z_j = h(a_j). \quad (5.48, 5.49)$$

$$\text{For each error-term:} \quad \frac{\partial E_n}{\partial w_{ji}} = \underbrace{\frac{\partial E_n}{\partial a_j}}_{\equiv \delta_j} \frac{\partial a_j}{\partial w_{ji}} \quad (5.50)$$

$$\text{So, from 5.48:} \quad \frac{\partial E_n}{\partial w_{ji}} = \delta_j z_i \quad (5.53)$$

$$\text{In the network:} \quad \delta_j \equiv \frac{\partial E_n}{\partial a_j} = \sum_k \frac{\partial E_n}{\partial a_k} \frac{\partial a_k}{\partial a_j} \quad \text{where } j \rightarrow \{k\} \quad (5.55)$$

$$\text{Algorithm:} \quad \delta_j = h'(a_j) \sum_k w_{kj} \delta_k \quad \text{as} \quad \frac{\partial a_k}{\partial a_j} = \frac{\partial a_k}{\partial z_j} \frac{\partial z_j}{\partial a_j} \quad (5.56)$$

Backpropagation Algorithm

The formula for the update of a given unit depends only on the 'later' (i.e. closer to the output) layers:




Figure: 5.7

Hence the backpropagation algorithm is:

- ▶ Apply input \mathbf{x} , and **forward propagate** to find the hidden and output activations.
- ▶ Evaluate δ_k directly for the output units.
- ▶ **Back propagate** the δ 's to obtain a δ_j for each hidden unit.
- ▶ Evaluate the derivatives $\frac{\partial E_n}{\partial w_{ji}} = \delta_j z_i$.

Computational Efficiency

The back-propagation algorithm is computationally more efficient than standard numerical minimization of E_n . Suppose that W is the total number of weights and biases in the network.

- ▶ **Backpropagation**: The evaluation is $O(W)$ for large W , as there are many more weights than units.
- ▶ **Standard approach**: Perturb each weight, and forward propagate to compute the change in E_n . This requires $W \times O(W)$ computations, so the total complexity is $O(W^2)$.

Jacobian Matrix

The properties of the network can be investigated via the [Jacobian](#)

$$J_{ki} = \frac{\partial y_k}{\partial x_i} \quad (5.70)$$

For example, (small) errors can be propagated through the trained network:

$$\Delta y_k \simeq \frac{\partial y_k}{\partial x_i} \Delta x_i \quad (5.72)$$

This is useful, but costly, as J_{ki} itself depends on \mathbf{x} . However, note that

$$\frac{\partial y_k}{\partial x_i} = \sum_j \frac{\partial y_k}{\partial a_j} \frac{\partial a_j}{\partial x_i} = \sum_j w_{ji} \frac{\partial y_k}{\partial a_j} \quad (5.74)$$

The required derivatives $\partial y_k / \partial a_j$ can be efficiently computed by [backpropagation](#).